

PCML Project 1

Yoam Hazot, Maximilian Mordig, Timon Zimmermann
EPFL Lausanne

Abstract—In this report, we report our results applying machine learning algorithms to a CERN dataset to detect Higgs bosons. We have a validation list which allows us to calibrate our models. We apply both linear regression with polynomial features (including cross-features) and logistic regression. For both, we use regularization and perform cross-validation over the regularization parameter λ , the degrees (for non-cross features), the degrees for cross-features, averaging all over a number of seeds used to select the training and the testing set. The best score we achieved was 81.153% with linear regression. Understanding the physics behind the measurement may definitely be a key to improve upon the results.

I. INTRODUCTION

The provided training dataset consists of 250000 data rows, each row contains 30 features and a prediction whether a particular confirms or rejects the presence of the Higgs boson. In the following, we explain how we obtained the best score of 81.153%.

A. Data Exploration

Before trying to fit a linear regression or similar to the data, it is important to get a rough idea of how it looks like.

We start by looking at the number of NaNs per column. We see that there are about seven features with 175000 NaNs values out of the total 250000 NaNs, as seen in figure 1. In the end, we choose to remove no features (based on too many NaN values) because this deteriorates the results. It seems that they contain useful information as the figure shows (comparing the data where the Higgs boson is detected and where it is not detected). In the code, we chose not to loop over this as well because the code will take too much time (do it yourself if you wish).

Instead of removing the features with too many NaN values, we choose to replace them by the median of the values of the feature. The median is more robust than the mean. This improves over the method where no preprocessing at all is applied.

The interesting thing is that roughly 30% of the provided data has a predicted Higgs boson, hence the trivial estimator can just assign "not predicted" to each data row which gives an overall prediction success of 70%. Hence, the benchline is to perform better than 70% of correct predictions. With no preprocessing of the data at all (no replacement or removal of the data with NaN values), just standardizing the data to have zero mean and unit variance, we achieved 74% of success (according to the Kaggle platform).

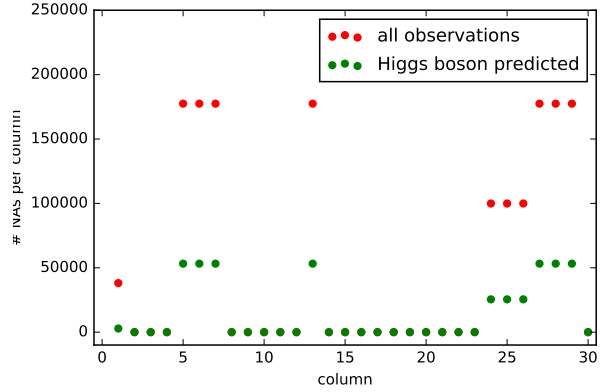


Figure 1. Number of NaN values in the data for each feature.

B. Applying PCA

Because the model becomes quite complex as the number of features increases, we implemented PCA (principal component analysis) to reduce the number of features. Choosing to keep a variance of 0.9999 reduces the 30 features to 15 features. We further plotted histograms for each of the features that we kept to check they do not exhibit any odd behaviour or if anything is special that distinguishes data with the Higgs boson from the data without the Higgs boson. We also plotted the cross-dependence for each pair of features, but it is difficult for us to extract the features from the plots in an efficient way. In other words, we did not use these plots except for gross error checking. We did not notice anything special. See the ipynb notebook file for the plots or the "Images" directory in the same folder.

C. Adding Polynomial Features Including Cross-Features

We consequently added polynomial non-cross-related features and did cross-validation on them. But we thought that given the correlation figures we had produced before, there seemed some significant correlation. Hence, we programmed a method to also generate cross-features. For instance, for features $\{x, y\}$, we generate the features $\{1, x, y, x^2, xy, y^2, x^3, x^2y, xy^2, y^3\}$ up to degree 3. We explicitly avoid duplicates xy and yx , which would both be created in a naive approach, by relying on the unique factorization of primes (our own idea, see the code for this) to generate a mask that picks exactly one among $\{xy, yx\}$. In the program, we generate the cross-related features from

degree 0 to some degree d_{cross} and then continue to generate the non-cross-related features from degree $d_{cross} + 1$ to d . In a better approach, to avoid computation complexity and overfitting, one could bring in the patterns observed from the correlations graphs by eye, but this limits the systematic applicability of the model.

We are aware of the fact that the cross-related features introduce a lot more complexity to the model. To avoid overfitting, we rely on the cross-validation rejecting the overfitted models. Practically, we split the provided data into 5% of training data and 95% of testing data, which should be enough for detecting overfitted models.

D. Using Cross-Validation to Select the Best Model

As mentioned above, the most important is the model selection among the class of possible models using k-cross-validation with $k = 5$. In fact, the list of models has become quite large. Each generated model depends on the cross-feature degree d_{cross} , the polynomial degree d , the regularization parameter λ , which features to remove from the very start as well as the variance to keep for the PCA dimensionality reduction. Additionally, there is another loop to average over the seeds used to select the testing and training sets for cross-validation. We were forced to decide to remove some complexity.

Changing one parameter at a time (*ceteris paribus*) and running cross-validation, we decided not to remove any features with too many NaN values and to keep 0.9999 of the variance (total of 1) with PCA. With the removed NaN values, we achieved the best testing error estimates on the provided dataset, but scored worse on Kaggle. We interpret the removal of features as being too biased towards the provided data sample because we determined the NaN columns also using the testing data (of the provided training data). What concerns the PCA, we were forced to reduce the dimensionality, so we thought, keeping 15 features with a retained variance of 0.9999 is still fine. For the seeds, we noticed that the training and testing errors do not change very much, so we restricted to two seeds (this is risky, but it takes too much time otherwise).

We also fixed the cross-feature degree to $d_{cross} = 2$, achieving best overall cross-validation results among all cross-features degrees ($d_{cross} \leq 3$ due to memory). We checked this manually instead of putting a loop around it. Finally, in the code, we have the following loops: over the seeds, over the (non-cross-feature) degrees and over λ (to select the best λ for a given degree).

Due to the polynomial (cross and non-cross) features, our model has a lot of variables in the end and the prediction of the test file ('test.csv') for the Kaggle platform may encounter memory errors. That is why we are predicting and writing the data in chunks.

Finally, we always used the exact method to compute the weights (involving matrix inversion). We found that the




Figure 2. Prediction success rate with degree of polynomial with penalised linear regression. Only curves for two seeds are plotted, but they reflect the general behaviour.

direct method was faster than (stochastic) gradient descent and more reliable.

For the cross-validation, we question whether it is good to select the method with least error or rather the method with best prediction accuracy (still using the least squares function for computing the weights). For this, see figure 2. For linear regression, it does not make sense to plot the least squares error as discussed in the lectures. We consequently looked at the number of correct predictions and plotted it as a function of the degree, where the weights were still obtained with the least squares function. We see that the testing error is slightly worse than the training error as expected and that we should best use degree 8. Here, we used cross-features up to degree 2.

E. Applying Logistic Regression

We also tried to apply (penalised) logistic regression, but we did not get very good results with it. This time, there was no exact formula available and we had to use stochastic gradient descent, which we checked was better than pure gradient descent. We could not find any good parameters for the method to converge for any degree and seed. For linear regression, we consequently used the normal equations for the exact solution, but this is not possible here. As a side remark, for computational purposes, we rewrote the involved sigmoid function $\sigma(x) = \frac{e^{-x}}{1+e^{-x}} = \frac{1}{1+e^x}$ and/or approximating $\ln(1 + e^x) \approx x$ for large $x \geq 10$.

F. Possibility for Improvements and Outlook

We obtained 81% success with our method presented here, which is not that bad. However, the machine learning approach to select the features appears very brutal to us. It was definitely a good idea to introduce cross-feature features, but greater care should be taken, which features to combine as this might significantly reduce the model complexity. We had another implementation of logistic regression (but forgot to

commit it), where we achieved 80%, which appears strange to us because logistic regression as a classification method should outperform linear regression.

With the polynomial features, we are for instance not able to capture a more complex relationship like an exponential or a logarithm. It would definitely prove useful to understand the physical significance of the provided measurements and develop an approximate parametrized physical model which we then fit with machine learning. For instance, we learned in the end that the "JET" parameter was a categorical variable and that it yielded much better results if we built a new model for each value of this variable (but we did not have enough time to include it). In the end, we are rather not content with 81% because we would at least need 99% prediction accuracy because other indirect methods of detecting the Higgs boson are scarce.